

Retrieval-Augmented Generation (RAG)

Retrieval-Augmented Generation (RAG) is a technique that combines information retrieval with a text generator model.

RAG systems first retrieve relevant documents from a knowledge base (often using vector databases like ChromaDB or pgvector) based on a user's query.

The retrieved context is then passed to a Large Language Model (LLM) alongside the original query.

This allows the LLM to generate answers that are more accurate, up-to-date, and grounded in specific facts, reducing hallucinations.